
Data Mining

Tamás Budavári - budavari@jhu.edu

Class 5

- Regularization
 - Principal Component Analysis
 - Lagrange multipliers
 - Explained variance
-

Linear Regression

- A linear combination of known $\phi_k(\cdot)$ **basis** functions

$$f(t; \boldsymbol{\beta}) = \sum_{k=1}^K \beta_k \phi_k(t)$$

It's a dot product with $\boldsymbol{\beta} = (\beta_1, \dots, \beta_K)^T$

- Evaluated at all data points $x = (x_1, x_2, \dots, x_N)$

$$f(x; \boldsymbol{\beta}) = X\boldsymbol{\beta}$$

where $X_{ik} = \phi_k(x_i)$

Method of Least Squares

- At the optimum

$$\hat{\beta} = (X^T X)^{-1} X^T y$$


- Hat matrix

$$\hat{y} = X\hat{\beta} = Hy$$

In [1]: `%pylab inline`

Populating the interactive namespace from numpy and matplotlib

```
In [2]: # generate a dataset with errors  
x = 3 * random.rand(50) # uniform between 0 and 3  
eps = 1 * random.randn(x.size) # normal noise  
y = 10*cos(x+1) + eps; plot(x,y,'bx');
```




```

In [3]: def poly(x,n):
        X = np.zeros((x.size,n+1));
        for i in range(X.shape[1]):
            X[:,i] = x**i
        return X

        # show data in black
        plot(x,y,'kx'); ylim(-20,20);

        xx = np.linspace(-1,4,500) # grid on x
        color = 'yrgbm' * 5 # color sequence
        for n in range(2,9):
            X = poly(x,n) # design matrix for fitting
            bHat = linalg.pinv(X).dot(y)
            yy = poly(xx,n).dot(bHat) # prediction
            plot(xx,yy,'-',c=color[n]);

```



Regularization

Penalize large coefficients in β

- **Ridge regression** uses L_2

$$\hat{\beta} = \operatorname{argmin}_{\beta} |y - X\beta|_2^2 + \lambda |\beta|_2^2$$

or even with a constant matrix Γ

$$\hat{\beta} = \operatorname{argmin}_{\beta} |y - X\beta|_2^2 + \lambda |\Gamma\beta|_2^2$$

- **Lasso regression** uses L_1

$$\hat{\beta} = \operatorname{argmin}_{\beta} |y - X\beta|_2^2 + \lambda |\beta|_1$$

L_1 yields sparse results

Different geometric meanings!

Linear Combinations

- Coefficients mix a given set of basis vectors, functions, images, shapes, ...

$$f(x; \beta) = \sum_k \beta_k \phi_k(x)$$


Fourier series



Discrete Cosine Transform (JPEG)



Spherical Harmonics



- What is a good basis like?

Principal Component Analysis

Statistical Learning

	Supervised	Unsupervised
Discrete	Classification	Clustering
Continuous	Regression	Dimensionality Reduction



Directions of Maximum Variance

- Let $X \in \mathbb{R}^N$ be a continuous random variable with $\mathbb{E}[X] = 0$ mean and covariance matrix C . What is the direction of maximum variance?

For any vector $a \in \mathbb{R}^N$

$$\text{Var}[a^T X] = \mathbb{E}[(a^T X)(X^T a)] = \mathbb{E}[a^T (XX^T) a]$$

so

$$\text{Var}[a^T X] = a^T \mathbb{E}[XX^T] a = a^T C a$$

We have to maximize this such that $a^2 = 1$

Constrained Optimization

- **Lagrange multiplier:** extra term with new parameter λ

$$\hat{a} = \arg \max_{a \in \mathbb{R}^N} [a^T C a - \lambda (a^2 - 1)]$$

- Partial derivatives vanish at optimum

$$\frac{\partial}{\partial \lambda} \rightarrow \hat{a}^2 - 1 = 0 \quad (\text{duh!})$$

$$\frac{\partial}{\partial a_k} \rightarrow ?$$

With indices

$$\max_{a \in \mathbb{R}^N} \left[\sum_{i,j} a_i C_{ij} a_j - \lambda \left(\sum_i a_i^2 - 1 \right) \right]$$

- Partial derivatives $\partial/\partial a_k$ vanish at optimum

$$\begin{aligned} \sum_{i,j} \frac{\partial a_i}{\partial a_k} C_{ij} a_j + \sum_{i,j} a_i C_{ij} \frac{\partial a_j}{\partial a_k} - 2\lambda \left(\sum_i a_i \frac{\partial a_i}{\partial a_k} \right) &= \\ &= \sum_{i,j} \delta_{ik} C_{ij} a_j + \sum_{i,j} a_i C_{ij} \delta_{jk} - 2\lambda \left(\sum_i a_i \delta_{ik} \right) = \\ &= \sum_j C_{kj} a_j + \sum_i a_i C_{ik} - 2\lambda a_k \end{aligned}$$

And back again...

- With vectors and matrices

$$C\hat{a} + C^T\hat{a} - 2\lambda\hat{a} = 0$$

but C is symmetric

$$C\hat{a} = \lambda\hat{a}$$

- Eigenproblem !!

Result

- The value of maximum variance is

$$\hat{a}^T C \hat{a} = \hat{a}^T \lambda \hat{a} = \lambda \hat{a}^T \hat{a} = \lambda$$

the largest eigenvalue λ_1

- The direction of maximum variance is the corresponding eigenvector a_1

$$Ca_1 = \lambda_1 a_1$$

- This is the **1st Principal Component**

2nd Principal Component

- Direction of largest variance uncorrelated to 1st PC

$$\hat{a} = \arg \max_{a \in \mathbb{R}^N} [a^T C a - \lambda (a^2 - 1) - \lambda' (a^T C a_1)]$$

- Partial derivatives vanish at optimum

$$2C\hat{a} - 2\lambda\hat{a} - \lambda'Ca_1 = 0$$

Result

- Multiply by a_1^T .

$$\begin{aligned} 2a_1^T C\hat{a} - 2a_1^T \lambda\hat{a} - a_1^T \lambda'Ca_1 &= 0 \\ 0 - 0 - \lambda'\lambda_1 &= 0 \rightarrow \lambda' = 0 \end{aligned}$$

- Still just an eigenproblem

$$C\hat{a} = \lambda\hat{a}$$

- Solution λ_2 and a_2

PCA

- Spectral decomposition or eigenvalue decomposition or eigendecomposition

Let $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_N \geq 0$ be the eigenvalues of C and e_1, \dots, e_N the corresponding eigenvectors

$$C = \sum_{k=1}^N \lambda_k (e_k e_k^T)$$

Consider $C e_l = \sum_k \lambda_k e_k (e_k^T e_l) = \lambda_l e_l$ for any l

- Matrix form

With diagonal Λ matrix of the eigenvalues and an E matrix of $[e_1, \dots, e_N]$

$$C = E \Lambda E^T$$

- The eigenvectors of largest eigenvalues capture the most variance

If keeping only $K < N$ eigenvectors, the best approximation is taking the first K PCs

$$C \approx \sum_{k=1}^K \lambda_k (e_k e_k^T) = E_K \Lambda_K E_K^T$$

New Coordinate System

- The E matrix of eigenvectors is a rotation, $E E^T = I$

$$Z = E^T X$$

- A truncated set of eigenvectors E_K defines a projection

$$Z_K = E_K^T X$$

and

$$X_K = E_K Z_K = E_K E_K^T X = P_K X$$

Detour: Projections

- If the square of a matrix is equal to itself

$$P^2 = P$$

- For example, projecting on the e unit vector

Scalar times vector

$$r' = e (e^T r) = e \beta_r$$

Or projection of vector r

$$r' = (e e^T) r = P r$$



Again

- The eigenvectors of largest eigenvalues capture the most variance

$$C \approx C_K = \sum_{k=1}^K \lambda_k (e_k e_k^T) = \sum_{k=1}^K \lambda_k P_k$$

- And the remaining eigenvectors span the subspace with the least variance

$$C - C_K = \sum_{l=K+1}^N \lambda_l P_l$$

Samples

- Set of N -vectors arranged in matrix $X = [x_1, x_2, \dots, x_n]$ with average of 0
This is NOT the random variable we talked about previously but the data matrix!

Sample covariance matrix is

$$C = \frac{1}{n-1} XX^T = \frac{1}{n-1} \sum_i x_i x_i^T$$

- Singular Value Decomposition (SVD)

$$X = UWV^T$$

where $U^T U = I$, W is diagonal, and $V^T V = I$

- Hence

$$C = \frac{1}{n-1} UWV^T VWU^T = \frac{1}{n-1} UW^2 U^T$$

So, if $C = E\Lambda E^T$ then $E = U$ and $\Lambda = \frac{1}{n-1} W^2$


Random Sample from Bivariate Normal

- See previous lecture

```
In [4]: from scipy.stats import norm
# generate multiple 2-D (column) vectors
S = norm.rvs(0,1,(2,20))
S[0,:] *= 4 # scale axis 0
f = +pi/4 # rotate by 45 degrees
R = array([[cos(f), -sin(f)],
           [sin(f),  cos(f)]])
X = R.dot(S)
X += np.array([[1],[3]]) # shift

figure(figsize=(5,5)); xlim(-15,15); ylim(-15,15);
plot(X[0,:],X[1:], 'o', alpha=0.9)
```

Out[4]: [



```
In [5]: # subtract sample mean
avg = mean(X, axis=1).reshape(X[:,1].size,1)
X -= avg
# sample covariance matrix
C = X.dot(X.T) / (X[0,:].size-1)
print ("Average\n", avg)
print ("Covariance\n", C)
```

```
Average
[[ 2.13266854]
 [ 3.81741732]]
Covariance
[[ 6.56960563  3.80359671]
 [ 3.80359671  3.21387998]]
```

```
In [6]: L, E = np.linalg.eig(C)
E, L
```

```
Out[6]: (array([[ 0.83773535, -0.54607644],
                [ 0.54607644,  0.83773535]]), array([ 9.04897404,  0.73451157]))
```

```
In [7]: E, L, E_same = np.linalg.svd(C)
E, L
```

```
Out[7]: (array([[ -0.83773535, -0.54607644],
                [-0.54607644,  0.83773535]]), array([ 9.04897404,  0.73451157]))
```

```
In [8]: E.dot(E.T)
```

```
Out[8]: array([[ 1.00000000e+00,  5.55111512e-17],
                [ 5.55111512e-17,  1.00000000e+00]])
```

```
In [9]: np.allclose( E.T, np.linalg.inv(E) )
```

```
Out[9]: True
```

```
In [10]: U, W, V = np.linalg.svd(X)
U, W**2 / (X[0,:].size-1)
```

```
Out[10]: (array([[ -0.83773535, -0.54607644],
                [-0.54607644,  0.83773535]]), array([ 9.04897404,  0.73451157]))
```



```
In [11]: # alternatively
         U, W**2 / (X.shape[1]-1)
```

```
Out[11]: (array([[ -0.83773535, -0.54607644],
                [ -0.54607644,  0.83773535]]), array([ 9.04897404,  0.73451157]))
```

```
In [12]: [ np.allclose( U.dot(U.T), np.eye(U.shape[0]) ),
         np.allclose( V.dot(V.T), np.eye(V.shape[0]) ) ]
```

```
Out[12]: [True, True]
```

```
In [13]: from sklearn import decomposition
         pca = decomposition.PCA(n_components=X.shape[0])
         pca.fit(X.T) # different convention: row vs col !!!
         pca.components_.T, pca.explained_variance_
```

```
Out[13]: (array([[ 0.83773535, -0.54607644],
                [ 0.54607644,  0.83773535]]), array([ 9.04897404,  0.73451157]))
```